

# Margo and Thallium



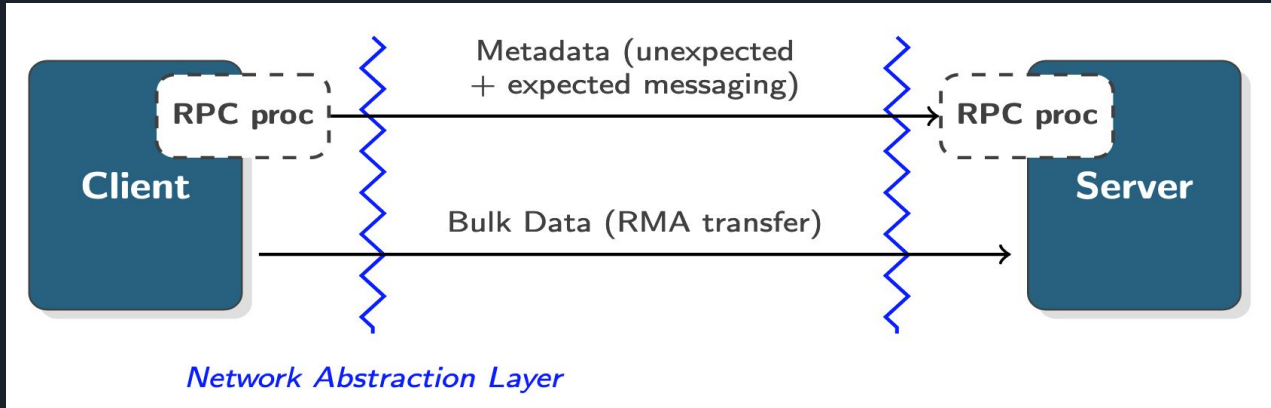
Argonne  
NATIONAL  
LABORATORY

Mochi Bootcamp  
February 6, 2020

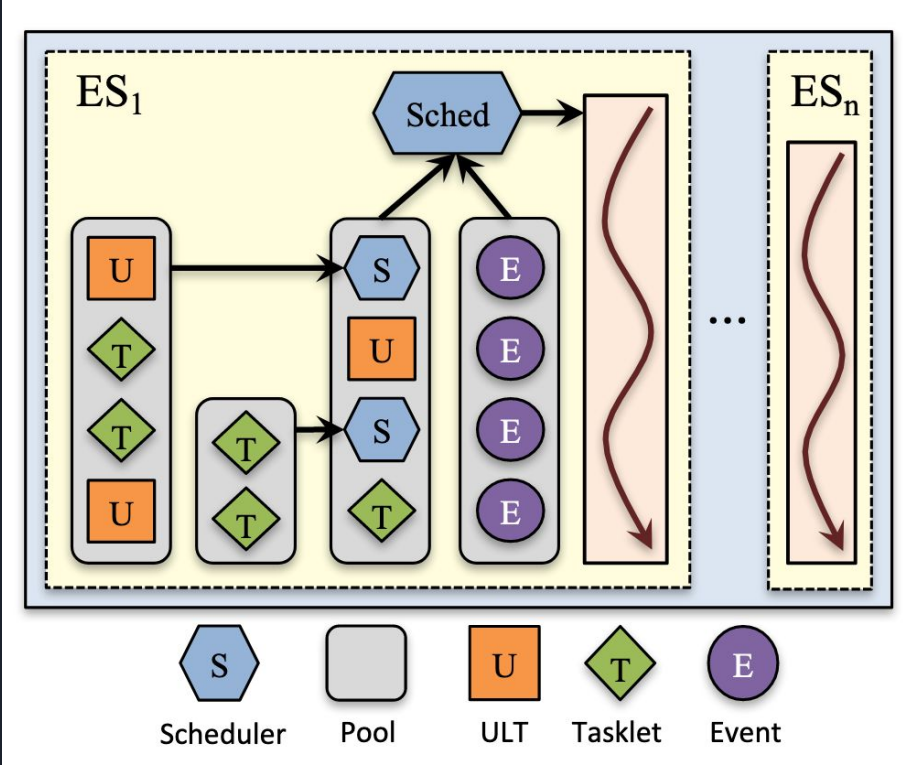
# A little bit about Mercury and Argobots



# Mercury: RPC and RDMA for HPC



- Multiple network abstractions
- Data serialization using macros
- Bulk data transfers (using RDMA when available)
- Explicit progress loop
- Callback-driven programming model



## Important concepts

- Execution streams
- User level threads
- Pools
- Schedulers

Argobots+Mercury = Margo

- Mercury progress placed in a ULT
- More natural programming model based on ULTs rather than callbacks



# Getting started with Margo



# Where to find the material

- <https://mochi.readthedocs.io>
- <https://xgitlab.cels.anl.gov/sds/mochi-doc>

```
git clone https://xgitlab.cels.anl.gov/sds/mochi-doc.git  
cd mochi-doc/code/margo
```

# Margo initialization



# Server initialization

```
#include <margo.h>

int main(int argc, char** argv)
{
    margo_instance_id mid = margo_init("tcp", MARGO_SERVER_MODE, 0, -1);
    assert(mid);

    hg_addr_t my_address;
    margo_addr_self(mid, &my_address);
    char addr_str[128];
    size_t addr_str_size = 128;
    margo_addr_to_string(mid, addr_str, &addr_str_size, my_address);
    margo_addr_free(mid, my_address);
    printf("Server running at address %s\n", addr_str);

    margo_wait_for_finalize(mid);

    return 0;
}
```



# Server initialization

```
#include <margo.h>
```

```
int main(int argc, char** argv)
{
```

```
    margo_instance_id mid = margo_init("tcp", MARGO_SERVER_MODE, 0, -1);
    assert(mid);
```

```
    hg_addr_t addr;
    margo_addr_t my_addr;
    char addr_str[128];
    size_t addr_str_size;
```

```
    margo_addr_to_string(mid, addr_str, &addr_str_size, my_address);
    margo_addr_free(mid, my_address);
    printf("Server running at address %s\n", addr_str);
```

```
    margo_wait_for_finalize(mid);
```

```
    return 0;
```

```
}
```

How many execution streams for RPC handlers?

Which protocol should we use?  
(tcp, na+sm, ofi+gni, etc.)

Should we run the progress loop  
in a separate execution stream?

Wait here until another thread  
calls `margo_finalize(mid)`

# Client initialization

```
#include <margo.h>

int main(int argc, char** argv)
{
    margo_instance_id mid = margo_init("tcp", MARGO_CLIENT_MODE, 0, 0);
    assert(mid);

    margo_finalize(mid);

    return 0;
}
```

Should we run the progress loop  
in a separate execution stream?

# “Hello World” RPC



# Simple “Hello World” RPC

```
#include <margo.h>

static const int TOTAL_RPCS = 4;
static int num_rpcs = 0;

static void hello_world(hg_handle_t h);
DECLARE_MARGO_RPC_HANDLER(hello_world)

int main(int argc, char** argv)
{
    ...
    hg_id_t rpc_id = MARGO_REGISTER(mid, "hello", void, void, hello_world);
    margo_registered_disable_response(mid, rpc_id, HG_TRUE);

    margo_wait_for_finalize(mid);
    return 0;
}
```

# Simple “Hello World” RPC

Server

All RPC handlers have this signature

```
static const int TO_CLIENT = 4;  
static int num_rpcs = 0;  
  
static void hello_world(hg_handle_t h);  
DECLARE_MARGO_RPC_HANDLER(hello_world)
```

Use this macro to declare the RPC handler

```
int main(int argc, char** argv)  
{
```

Register your RPC handler

```
    ...  
    hg_id_t rpc_id = MARGO_REGISTER(mid, "hello", void, void, hello_world);  
    margo_registered_disable_response(mid, rpc_id, HG_TRUE);
```

```
    margo_wait_for_finalize(mid);  
    return 0;  
}
```

This RPC doesn't send a response back to the client

# Simple “Hello World” RPC (cont’d)

```
static void hello_world(hg_handle_t h)
{
    hg_return_t ret;
    margo_instance_id mid = margo_hg_handle_get_instance(h);
    printf("Hello World!\n");
    num_rpcs += 1;

    ret = margo_destroy(h);
    assert(ret == HG_SUCCESS);

    if(num_rpcs == TOTAL_RPCCS) {
        margo_finalize(mid);
    }
}
DEFINE_MARGO_RPC_HANDLER(hello_world)
```

Don't forget to destroy the handle

The margo instance will be finalized after this RPC finishes, and `margo_wait_for_finalize` in main will return

Define the RPC handler

# Simple “Hello World” RPC (cont’d)

```
#include <margo.h>

int main(int argc, char** argv)
{
    ...
    hg_id_t hello_rpc_id = MARGO_REGISTER(mid, "hello", void, void, NULL);
    margo_registered_disable_response(mid, hello_rpc_id, HG_TRUE);

    hg_addr_t svr_addr;
    ret = margo_addr_lookup(mid, "tcp://localhost:1234", &svr_addr);

    hg_handle_t handle;
    ret = margo_create(mid, svr_addr, hello_rpc_id, &handle);
    ret = margo_forward(handle, NULL);
    ret = margo_destroy(handle);

    ret = margo_addr_free(mid, svr_addr);
    ...
}
```

# Simple “Hello World” RPC (cont’d)

```
#include <margo.h>
```

```
int main(int  
{
```

```
...
```

```
hg_id_t hello_rpc_id = MARGO_REGISTER(mid, "hello", void, void, NULL);
```

```
margo_registered_disable_response(mid, hello_rpc_id, HG_TRUE);
```

```
hg_addr_t svr_addr;
```

```
ret = margo_addr_lookup(mid, "tcp://localhost:1234", &svr_addr);
```

```
hg_handle_t handle;
```

```
ret = margo_create(mid, svr_addr, hello_rpc_id, &handle);
```

```
ret = margo_forward(handle, NULL);
```

```
ret = margo_destroy(handle);
```

```
ret = margo_addr_free(mid, svr_addr);
```

```
...
```

```
}
```

This RPC doesn't expect a response from the server

NULL used instead of function

Lookup the server's address

Create, forward (send), and destroy the RPC



Sending arguments,  
returning values



# Sending arguments, returning values

```
#ifndef PARAM_H  
#define PARAM_H
```

Include the Mercury macros

```
#include <mercury.h>  
#include <mercury_macros.h>
```

```
/* We use the Mercury macros to define the input  
 * and output structures along with the serialization functions.  
 */
```

```
MERCURY_GEN_PROC(sum_in_t,  
    ((int32_t)(x))\  
    ((int32_t)(y)))
```

Generate the definition of the  
sum\_in\_t structure and its  
serialization functions

```
MERCURY_GEN_PROC(sum_out_t, ((int32_t)(ret)))
```

```
#endif
```

# Sending arguments, returning values (cont'd)

```
#include <margo.h>
#include "types.h"

static void sum(hg_handle_t h);
DECLARE_MARGO_RPC_HANDLER(sum)

int main(int argc, char** argv) {
    ...
    hg_id_t rpc_id = MARGO_REGISTER(mid, "sum", sum_in_t, sum_out_t, sum);
    ...
}
```

Use the newly defined types  
when registering the RPC

# Sending arguments, returning values (cont'd)

```
static void sum(hg_handle_t h)
{
    hg_return_t ret;

    sum_in_t in;
    sum_out_t out;

    ret = margo_get_input(h, &in);

    out.ret = in.x + in.y;
    printf("Computed %d + %d = %d\n", in.x, in.y, out.ret);

    ret = margo_respond(h, &out);
    ret = margo_free_input(h, &in);
    ret = margo_destroy(h);
}
DEFINE_MARGO_RPC_HANDLER(sum)
```

Read the input sent by the client

Send a response with the output

Don't forget to free the input!

# Sending arguments, returning values (cont'd)

```
#include "types.h"
```

```
...
```

```
hg_id_t sum_rpc_id = MARGO_REGISTER(mid, "sum", sum_in_t, sum_out_t, NULL);
```

```
sum_in_t args;
```

```
args = { .x = 42, .y = 58 };
```

```
hg_handle_t h;
```

```
margo_create(mid, svr_addr, sum_rpc_id, &h);
```

```
margo_forward(h, &args);
```

Pass the input to forward

```
sum_out_t resp;
```

```
margo_get_output(h, &resp);
```

Read the RPC's output

```
printf("Got response: %d+%d = %d\n", args.x, args.y, resp.ret);
```

```
margo_free_output(h, &resp);
```

Free the RPC's output

```
margo_destroy(h);
```

```
...
```

# Attaching data to RPC handlers



# Attaching data to RPC handlers

```
typedef struct {  
    ...  
} mydata;  
  
...  
mydata* data = malloc(...);  
hg_id_t sum_rpc_id = MARGO_REGISTER(mid, "sum", sum_in_t, sum_out_t, NULL);  
margo_register_data(mid, sum_rpc_id, (void*)data, free);  
...  
  
const struct hg_info* info = margo_get_info(handle);  
mydata* data = (mydata*)margo_registered_data(mid, info->id);
```

Attach the pointer to the RPC

Give it the function to use to free the data (or NULL)

Inside the RPC handler, retrieve the pointer

# Bulk data transfers





# Bulk data transfers

```
#ifndef PARAM_H
#define PARAM_H

#include <mercury.h>
#include <mercury_macros.h>

MERCURY_GEN_PROC(sum_in_t,
    ((int32_t)(n))\
    ((hg_bulk_t)(bulk)))

MERCURY_GEN_PROC(sum_out_t, ((int32_t)(ret)))

#endif
```

We will send  $n$  integers using bulk data transfer

## Bulk data transfers (cont'd)

```
#include "types.h"
```

```
...
```

```
hg_id_t sum_rpc_id = MARGO_REGISTER(mid, "sum", sum_in_t, sum_out_t, NULL);
```

```
sum_in_t args;
```

```
int32_t values[10] = { 1,4,2,5,6,3,5,3,2,5 };
```

```
hg_size_t size = 10*sizeof(int32_t);
```

Expose the array by  
creating a bulk handle

```
hg_bulk_t local_bulk;
```

```
margo_bulk_create(mid, 1, (void**)&values, &size, HG_BULK_READ_ONLY, &local_bulk);
```

```
sum_in_t args = { .n = 10, .bulk = local_bulk };
```

```
...
```

```
margo_bulk_free(local_bulk);
```

Don't forget to free the bulk handle

```
...
```

# Bulk data transfers (cont'd)

```
#include "types.h"
```

```
static void sum(hg_handle_t h)
{
```

```
    hg_return_t ret;
    sum_in_t in;
    sum_out_t out;
    int32_t* values;
    hg_bulk_t local_bulk;
```

```
    margo_instance_id mid = margo_hg_handle_get_instance(h);
```

```
    const struct hg_info* info = margo_get_info(h);
    hg_addr_t client_addr = info->addr;
```

You can get the client's address this way

```
    ret = margo_get_input(h, &in);
```

```
    values = calloc(in.n, sizeof(*values));
    hg_size_t buf_size = in.n * sizeof(*values);
```

Allocate a local buffer to receive the values from the client

## Bulk data transfers (cont'd)

```
ret = margo_bulk_create(mid, 1, (void*)&values, &buf_size,  
                        HG_BULK_WRITE_ONLY, &local_bulk);  
  
ret = margo_bulk_transfer(mid, HG_BULK_PULL, client_addr,  
                          in.bulk, 0, local_bulk, 0, buf_size);  
  
out.ret = 0;  
int i;  
for(i = 0; i < in.n; i++) {  
    out.ret += values[i];  
}  
  
ret = margo_respond(h, &out);  
ret = margo_bulk_free(local_bulk);  
free(values);  
ret = margo_free_input(h, &in);  
ret = margo_destroy(h);  
}  
DEFINE_MARGO_RPC_HANDLER(sum)
```

Expose the local buffer

Pull data from the client

Free the bulk handle

# Non-blocking RPC



# Non-blocking RPC

```
sum_in_t args = { .x = 42, .y = 58 };
```

```
hg_handle_t h;
```

```
...
```

```
margo_request req;
```

```
margo_iforward(h, &args, &req);
```

Use `margo_iforward` to  
send and return immediately

```
printf("Waiting for reply...\n");
```

```
margo_wait(req);
```

Wait for the operation to  
complete (this will also free  
the request)

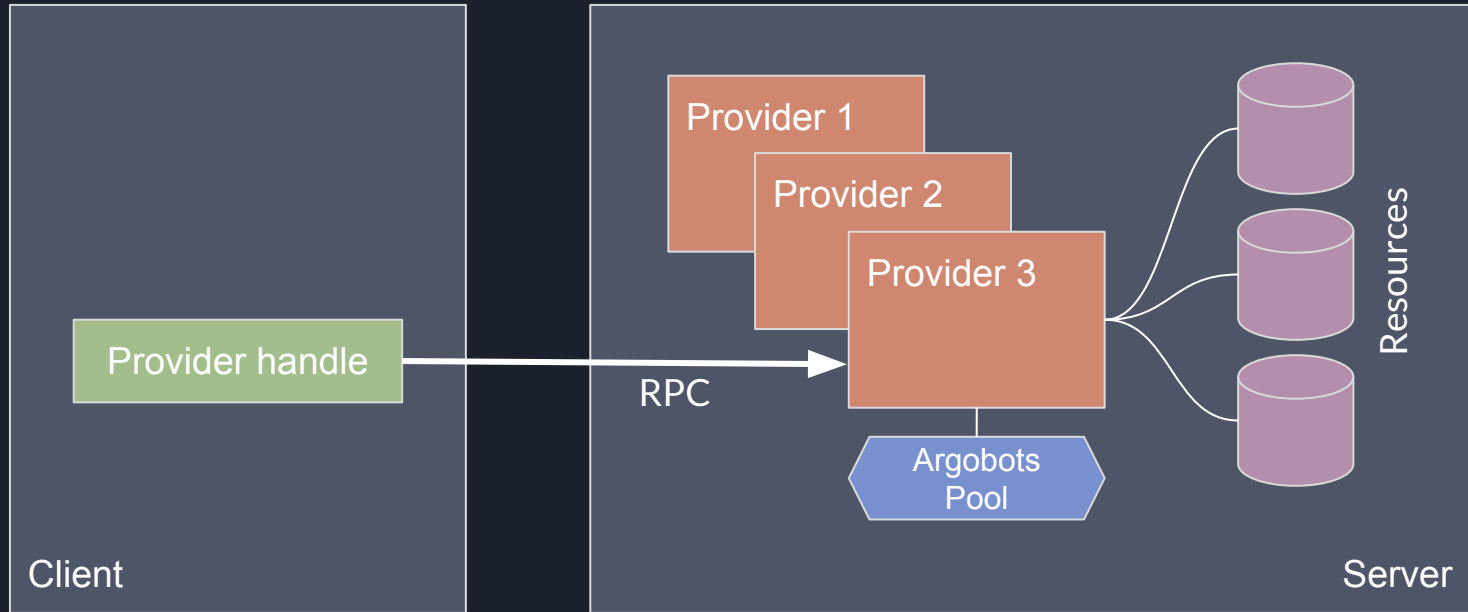
```
sum_out_t resp;
```

```
margo_get_output(h, &resp);
```

# Developing providers



# What is a provider?





# Developing providers

```
#ifndef PARAM_H
#define PARAM_H

#include <mercury.h>
#include <mercury_macros.h>

MERCURY_GEN_PROC(sum_in_t,
    ((int32_t)(x))\
    ((int32_t)(y)))

MERCURY_GEN_PROC(sum_out_t, ((int32_t)(ret)))

#endif
```



Just a reminder

## Developing providers (cont'd)

```
#ifndef __ALPHA_COMMON_H
#define __ALPHA_COMMON_H

#define ALPHA_SUCCESS  0
#define ALPHA_FAILURE -1

#endif
```

# Developing providers (cont'd)

```
typedef struct alpha_client* alpha_client_t;
#define ALPHA_CLIENT_NULL ((alpha_client_t)NULL)

typedef struct alpha_provider_handle *alpha_provider_handle_t;
#define ALPHA_PROVIDER_HANDLE_NULL ((alpha_provider_handle_t)NULL)

int alpha_client_init(margo_instance_id mid, alpha_client_t* client);

int alpha_client_finalize(alpha_client_t client);

int alpha_provider_handle_create(alpha_client_t client, hg_addr_t addr, uint16_t provider_id,
                                alpha_provider_handle_t* handle);

int alpha_provider_handle_ref_incr(alpha_provider_handle_t handle);

int alpha_provider_handle_release(alpha_provider_handle_t handle);

int alpha_compute_sum(alpha_provider_handle_t handle, int32_t x, int32_t y, int32_t* result);
```

# Developing providers (cont'd)

```
#include "types.h"
#include "alpha-client.h"

struct alpha_client {
    margo_instance_id mid;
    hg_id_t            sum_id;
    uint64_t           num_prov_hdl;
};

struct alpha_provider_handle {
    alpha_client_t client;
    hg_addr_t      addr;
    uint16_t       provider_id;
    uint64_t       refcount;
};
```

# Developing providers (cont'd)

```
int alpha_client_init(margo_instance_id mid, alpha_client_t* client)
{
    int ret = ALPHA_SUCCESS;
    alpha_client_t c = (alpha_client_t)calloc(1, sizeof(*c));
    if(!c) return ALPHA_FAILURE;
    c->mid = mid;
    hg_bool_t flag;
    hg_id_t id;
    margo_registered_name(mid, "alpha_sum", &id, &flag);
    if(flag == HG_TRUE) {
        margo_registered_name(mid, "alpha_sum", &c->sum_id, &flag);
    } else {
        c->sum_id = MARGO_REGISTER(mid, "alpha_sum", sum_in_t, sum_out_t, NULL);
    }
    *client = c;
    return ALPHA_SUCCESS;
}
```

# Developing providers (cont'd)

```
int alpha_client_init(margo_instance_id mid, alpha_client_t* client)
{
    int ret = ALPHA_SUCCESS;
    alpha_client_t c = (alpha_client_t){0};
    if(!c) return ALPHA_FAILURE;
    c->mid = mid;
    hg_bool_t flag;
    hg_id_t id;
    margo_registered_name(mid, "alpha_sum", &id, &flag);
    if(flag == HG_TRUE) {
        margo_registered_name(mid, "alpha_sum", &c->sum_id, &flag);
    } else {
        c->sum_id = MARGO_REGISTER(mid, "alpha_sum", sum_in_t, sum_out_t, NULL);
    }
    *client = c;
    return ALPHA_SUCCESS;
}
```

Check if the RPC is already registered

If it is, get the ID

Otherwise, register it

## Developing providers (cont'd)

```
int alpha_client_finalize(alpha_client_t client)
{
    if(client->num_prov_hdl != 0) {
        fprintf(stderr,
            "Warning: %d provider handles not released when alpha_client_finalize was called\n",
            client->num_prov_hdl);
    }
    free(client);
    return ALPHA_SUCCESS;
}
```

# Developing providers (cont'd)

```
int alpha_provider_handle_create(alpha_client_t client, hg_addr_t addr, uint16_t provider_id,
                                alpha_provider_handle_t* handle)
{
    if(client == ALPHA_CLIENT_NULL) return ALPHA_FAILURE;
    alpha_provider_handle_t ph = (alpha_provider_handle_t)calloc(1, sizeof(*ph));
    if(!ph) return ALPHA_FAILURE;
    hg_return_t ret = margo_addr_dup(client->mid, addr, &(ph->addr));
    if(ret != HG_SUCCESS) {
        free(ph);
        return ALPHA_FAILURE;
    }
    ph->client      = client;
    ph->provider_id = provider_id;
    ph->refcount     = 1;
    client->num_prov_hdl += 1;
    *handle = ph;
    return ALPHA_SUCCESS;
}
```

Duplicate the address, so the caller can safely free the original



# Developing providers (cont'd)

```
int alpha_provider_handle_ref_incr(alpha_provider_handle_t handle)
{
    handle->refcount += 1;
    return ALPHA_SUCCESS;
}

int alpha_provider_handle_release(alpha_provider_handle_t handle)
{
    handle->refcount -= 1;
    if(handle->refcount == 0) {
        margo_addr_free(handle->client->mid, handle->addr);
        handle->client->num_prov_hdl -= 1;
        free(handle);
    }
    return ALPHA_SUCCESS;
}
```

Free the provider handle's  
underlying address

## Developing providers (cont'd)

```
int alpha_compute_sum(alpha_provider_handle_t handle, int32_t x, int32_t y, int32_t* result)
{
    hg_handle_t    h;
    sum_in_t       in = { .x = x, .y = y };
    sum_out_t      out;
    hg_return_t     ret;
    ret = margo_create(handle->client->mid, handle->addr, handle->client->sum_id, &h);
    ret = margo_provider_forward(handle->provider_id, h, &in);
    ret = margo_get_output(h, &out);
    *result = out.ret;
    margo_free_output(h, &out);
    margo_destroy(h);
    return ALPHA_SUCCESS;
}
```

Use `margo_provider_forward`  
instead of `margo_forward`

# Developing providers (cont'd)

```
#define ALPHA_ABT_POOL_DEFAULT ABT_POOL_NULL

typedef struct alpha_provider* alpha_provider_t;
#define ALPHA_PROVIDER_NULL ((alpha_provider_t)NULL)
#define ALPHA_PROVIDER_IGNORE ((alpha_provider_t*)NULL)

int alpha_provider_register(
    margo_instance_id mid,
    uint16_t provider_id,
    ABT_pool pool,
    alpha_provider_t* provider);

int alpha_provider_destroy(
    alpha_provider_t provider);
```

# Developing providers (cont'd)

```
#include "alpha-server.h"
#include "types.h"

struct alpha_provider {
    margo_instance_id mid;
    hg_id_t sum_id;
    /* other provider-specific data */
};

static void alpha_finalize_provider(void* p);

DECLARE_MARGO_RPC_HANDLER(alpha_sum_ult);
static void alpha_sum_ult(hg_handle_t h);
/* add other RPC declarations here */
```

This function will be called either by `alpha_provider_destroy` or when the Margo instance is finalized

## Developing providers (cont'd)

```
int alpha_provider_register(margo_instance_id mid, uint16_t provider_id, ABT_pool pool,
                           alpha_provider_t* provider)
{
    alpha_provider_t p;
    hg_id_t id;
    hg_bool_t flag;

    flag = margo_is_listening(mid);
    if(flag == HG_FALSE) {
        fprintf(stderr, "alpha_provider_register(): margo instance is not a server.");
        return ALPHA_FAILURE;
    }

    margo_provider_registered_name(mid, "alpha_sum", provider_id, &id, &flag);
    if(flag == HG_TRUE) {
        fprintf(stderr, "a provider with the same provider id already exists.\n");
        return ALPHA_FAILURE;
    }
}
```

No point in creating a provider if it cannot receive anything

Check if RPC already registered for this provider id

## Developing providers (cont'd)

```
p = (alpha_provider_t)calloc(1, sizeof(*p));
if(p == NULL)
    return ALPHA_FAILURE;

p->mid = mid;

id = MARGO_REGISTER_PROVIDER(mid, "alpha_sum",
    sum_in_t, sum_out_t,
    alpha_sum_ult, provider_id, pool);
margo_register_data(mid, id, (void*)p, NULL);
p->sum_id = id;
/* add other RPC registration here */

margo_provider_push_finalize_callback(mid, p, &alpha_finalize_provider, p);

*provider = p;
return ALPHA_SUCCESS;
}
```

# Developing providers (cont'd)

```
p = (alpha_provider_t)calloc(1, sizeof(*p));
if(p == NULL)
    return ALPHA_FAILURE;

p->mid = mid;

id = MARGO_REGISTER_PROVIDER(mid, "alpha_sum",
    sum_in_t, sum_out_t,
    alpha_sum_ult, provider_id, pool);
margo_register_data(mid, id, (void*)p, NULL);
p->sum_id = id;
/* add other RPC registration here */

margo_provider_push_finalize_callback(mid, p, &alpha_finalize_provider, p);

*provider = p;
return ALPHA_SUCCESS;
}
```

Use `MARGO_REGISTER_PROVIDER`  
instead of `MARGO_REGISTER`

Pass an Argobots pool to run RPC

Attach a pointer to the  
provider to the RPC

Tell Margo to destroy the provider if  
the Margo instance is finalized

## Developing providers (cont'd)

```
static void alpha_finalize_provider(void* p)
{
    alpha_provider_t provider = (alpha_provider_t)p;
    margo_deregister(provider->mid, provider->sum_id);
    /* deregister other RPC ids ... */
    free(provider);
}

int alpha_provider_destroy(
    alpha_provider_t provider)
{
    /* pop the finalize callback */
    margo_provider_pop_finalize_callback(provider->mid, provider);
    /* call the callback */
    alpha_finalize_provider(provider);

    return ALPHA_SUCCESS;
}
```

Deregister the RPC

Tell Margo there is no need anymore to call the finalize callback



# Developing providers (cont'd)

```
static void alpha_sum_ult(hg_handle_t h)
{
    hg_return_t ret;
    sum_in_t    in;
    sum_out_t   out;

    margo_instance_id mid = margo_hg_handle_get_instance(h);
    const struct hg_info* info = margo_get_info(h);
    alpha_provider_t provider = (alpha_provider_t)margo_registered_data(mid, info->id);

    ret = margo_get_input(h, &in);

    out.ret = in.x + in.y;
    printf("Computed %d + %d = %d\n", in.x, in.y, out.ret);

    ret = margo_respond(h, &out);
    ret = margo_free_input(h, &in);
}
DEFINE_MARGO_RPC_HANDLER(alpha_sum_ult)
```

This is how to retrieve the pointer to the provider that we have attached earlier



# Getting started with Thallium



# Where to find the material

- <https://mochi.readthedocs.io>
- <https://xgitlab.cels.anl.gov/sds/mochi-doc>

```
git clone https://xgitlab.cels.anl.gov/sds/mochi-doc.git  
cd mochi-doc/code/thallium
```

# Initializing Thallium



# Initializing Thallium

```
#include <iostream>
#include <thallium.hpp>
```

```
namespace tl = thallium;
```

```
int main(int argc, char** argv) {
```

```
    tl::engine myEngine("tcp", THALLIUM_SERVER_MODE);
    std::cout << "Server running at address " << myEngine.self() << std::endl;
```

```
    return 0;
```

```
}
```

The engine is the core class running Mercury progress

You may add two more parameters:

- `bool use_progress_thread`
- `int num_rpc_threads`

The engine's destructor will block until finalized

# Initializing Thallium (cont'd)

```
#include <thallium.hpp>

namespace tl = thallium;

int main(int argc, char** argv) {

    tl::engine myEngine("tcp", THALLIUM_CLIENT_MODE);

    return 0;
}
```

When initialized as a client, the engine will finalize itself when its destructor is called, instead of blocking

# “Hello World” RPC



# “Hello world” RPC

```
#include <iostream>
#include <thallium.hpp>

namespace tl = thallium;

void hello(const tl::request& req) {
    std::cout << "Hello World!" << std::endl;
}

int main(int argc, char** argv) {

    tl::engine myEngine("tcp://127.0.0.1:1234", THALLIUM_SERVER_MODE);
    myEngine.define("hello", hello).disable_response();

    return 0;
}
```

RPC handlers must take a reference to a request object

Defines the “hello” RPC

Tell thallium this RPC doesn't send a response



## “Hello world” RPC (cont’d)

```
#include <thallium.hpp>
```

```
namespace tl = thallium;
```

```
int main(int argc, char** argv) {
```

```
    tl::engine myEngine("tcp", THALLIUM_CLIENT_MODE);
```

```
    tl::remote_procedure hello = myEngine.define("hello").disable_response();
```

```
    tl::endpoint server = myEngine.lookup("tcp://127.0.0.1:1234");
```

```
    hello.on(server)();
```

```
    return 0;
```

```
}
```

Define the RPC without  
the function pointer

Lookup the server's address

Call the RPC on the server

Sending arguments,  
returning values



# Sending arguments, returning values

```
#include <iostream>
#include <thallium.hpp>
```

```
namespace tl = thallium;
```

```
void sum(const tl::request& req, int x, int y) {
    std::cout << "Computing " << x << "+" << y << std::endl;
    req.respond(x+y);
}
```

RPC handlers may take additional arguments

And send a response

```
int main(int argc, char** argv) {

    tl::engine myEngine("tcp://127.0.0.1:1234", THALLIUM_SERVER_MODE);
    myEngine.define("sum", sum);

    return 0;
}
```

# Sending arguments, returning values (cont'd)

```
#include <iostream>
#include <thallium.hpp>

namespace tl = thallium;

int main(int argc, char** argv) {

    tl::engine myEngine("tcp", THALLIUM_CLIENT_MODE);
    tl::remote_procedure sum = myEngine.define("sum");
    tl::endpoint server = myEngine.lookup("tcp://127.0.0.1:1234");
    int ret = sum.on(server)(42,63);
    std::cout << "Server answered " << ret << std::endl;

    return 0;
}
```

# Sending arguments, returning values (cont'd)

Client

```
#include <iostream>
#include <thallium.hpp>
```

```
namespace tl = thallium;
```

```
int main(int argc, char** argv) {
```

```
    tl::engine myEngine("tcp", THALLIUM_CLIENT_MODE);
    tl::remote_procedure sum = myEngine.define("sum");
    tl::endpoint server = myEngine.lookup("tcp://127.0.0.1:1234");
    int ret = sum.on(server)(42,63);
    std::cout << "Server answered " << ret << std::endl;
```

```
    return 0;
```

```
}
```

We should now call the RPC  
with the extra arguments

And receive the result

**Warning: make sure  
your types match  
between RPC handlers  
and RPC calls!**

# Registering lambdas as RPC handlers



# Registering lambdas as RPC handlers

```
#include <iostream>
#include <thallium.hpp>

namespace tl = thallium;

int main(int argc, char** argv) {

    tl::engine myEngine("tcp://127.0.0.1:1234", THALLIUM_SERVER_MODE);

    std::function<void(const tl::request&, int, int)> sum =
        [](const tl::request& req, int x, int y) {
            std::cout << "Computing " << x << "+" << y << std::endl;
            req.respond(x+y);
        };

    myEngine.define("sum", sum);

    return 0;
}
```

# Non-blocking RPCs





# Non-blocking RPCs

```
tl::engine myEngine("tcp", THALLIUM_CLIENT_MODE);
tl::remote_procedure sum = myEngine.define("sum");
tl::endpoint server = myEngine.lookup("tcp://127.0.0.1:1234");
auto request = sum.on(server).async(42,63);
// do something else ...
// check if request completed
bool completed = request.received();
// ...
// actually wait on the request and get the result out of it
int ret = request.wait();
std::cout << "Server answered " << ret << std::endl;
```

# Non-blocking RPCs

```
tl::engine myEngine("tcp", THALLIUM_CLIENT_MODE);
tl::remote_procedure sum = myEngine.define("sum");
tl::endpoint server = myEngine.lookup("tcp://127.0.0.1:1234");
auto request = sum.on(server).async(42,63);
// do something else ...
// check if request completed
bool completed = request.received();
// ...
// actually wait on the request and get the result
int ret = request.wait();
std::cout << "Server answered " << ret << std::endl;
```

Non-blocking test for completion

wait() is required even if  
received() returned true

Properly finalizing a server



# Properly finalizing a server

```
#include <iostream>
#include <thallium.hpp>

namespace tl = thallium;

int main(int argc, char** argv) {

    tl::engine myEngine("tcp://127.0.0.1:1234", THALLIUM_SERVER_MODE);

    std::function<void(const tl::request&, int, int)> sum =
        [&myEngine](const tl::request& req, int x, int y) {
            std::cout << "Computing " << x << "+" << y << std::endl;
            req.respond(x+y);
            myEngine.finalize();
        };
    myEngine.define("sum", sum);

    return 0;
}
```

`finalize()` can be called  
from an RPC handler

# Sending/receiving STL data structures





# All the C++14 data structures are supported

- **Any basic type** (int, float,...)
- `std::array<T>`
- `std::complex<T>`
- `std::deque<T>`
- `std::forward_list<T>`
- `std::list<T>`
- `std::map<K,V>`
- `std::multimap<K,V>`
- `std::multiset<T>`
- `std::pair<U,V>`
- `std::set<T>`
- `std::string`
- `std::tuple<T...>`
- `std::unordered_map<K,V>`
- `std::unordered_multimap<K,V>`
- `std::unordered_multiset<T>`
- `std::unordered_set<T>`
- `std::vector<T>`

Any composition of those types can be serialized by Thallium as well, for example  
`std::vector<std::tuple<std::pair<int,double>,std::list<int>>>`

# Example: sending/receiving a `std::string`

```
#include <string>
#include <iostream>
#include <thallium.hpp>
#include <thallium/serialization/stl/string.hpp>
```

Don't forget to include the right header!

```
namespace tl = thallium;

void hello(const tl::request& req, const std::string& name) {
    std::cout << "Hello " << name << std::endl;
}

int main(int argc, char** argv) {

    tl::engine myEngine("tcp://127.0.0.1:1234", THALLIUM_SERVER_MODE);
    myEngine.define("hello", hello).disable_response();

    return 0;
}
```

# Example: sending/receiving a `std::string`

```
#include <string>
#include <thallium.hpp>
#include <thallium/serialization/stl/string.hpp>
```

Don't forget to include the right header!

```
namespace tl = thallium;
```

```
int main(int argc, char** argv) {
```

```
    tl::engine myEngine("tcp", THALLIUM_CLIENT_MODE);
    tl::remote_procedure hello = myEngine.define("hello").disable_response();
    tl::endpoint server = myEngine.lookup("tcp://127.0.0.1:1234");
    std::string name = "Matthieu";
    hello.on(server)(name);
```

```
    return 0;
```

```
}
```

`hello.on(server) ("Matthieu")` would not work because template type deduction would infer a `const char*` instead of an `std::string`



Sending/receiving  
instances of custom class





# Example: a point class

```
class point {  
  
    private:  
  
        double x, y, z;  
  
    public:  
  
        point(double a=0.0, double b=0.0, double c=0.0)  
        : x(a), y(b), z(c) {}  
  
        template<typename A>  
        void serialize(A& ar) {  
            ar & x;  
            ar & y;  
            ar & z;  
        }  
};
```



## Other serialization methods

- `template "serialize"` member function
- `template "save"` and `"load"` member functions
- `template "serialize"` external function
- `template "save"` and `"load"` external functions
- `"read"` and `"write"` methods on archive type to access raw data

See <https://mochi.readthedocs.io> for more information.

# Bulk data transfers



# Bulk data transfers

```
tl::engine myEngine("tcp", MARGO_CLIENT_MODE);  
tl::remote_procedure remote_do_rdma = myEngine.define("do_rdma").disable_response();  
tl::endpoint server_endpoint = myEngine.lookup("tcp://127.0.0.1:1234");
```

```
std::string buffer = "Matthieu";  
std::vector<std::pair<void*, std::size_t>> segments(1);  
segments[0].first = (void*)&buffer[0];  
segments[0].second = buffer.size();
```

Expose one or more segments by using a vector of <pointer, size> pairs

```
tl::bulk myBulk = myEngine.expose(segments, tl::bulk_mode::read_only);
```

```
remote_do_rdma.on(server_endpoint)(myBulk);
```

The server is only going to read from this memory

## Bulk data transfers (cont'd)

```
tl::engine myEngine("tcp://127.0.0.1:1234", THALLIUM_SERVER_MODE);

std::function<void(const tl::request&, tl::bulk&)> f =
    [&myEngine](const tl::request& req, tl::bulk& b) {
        tl::endpoint ep = req.get_endpoint();
        std::string s(6,0);
        std::vector<std::pair<void*,std::size_t>> segments(1);
        segments[0].first = (void*)(s.data());
        segments[0].second = s.size();
        tl::bulk local = myEngine.expose(segments, tl::bulk_mode::write_only);
        b.on(ep) >> local;
        std::cout << "Server received bulk: ";
        std::cout << s;
        std::cout << std::endl;
    };

myEngine.define("do_rdma",f).disable_response();
```

# Bulk data transfers (cont'd)

```
tl::engine myEngine("tcp://127.0.0.1:1234", THALLIUM_SERVER_MODE);

std::function<void(const tl::request&, tl::bulk&)> f =
    [&myEngine](const tl::request& req, tl::bulk& b) {
        tl::endpoint ep = req.get_endpoint();
        std::string s(6,0);
        std::vector<std::pair<void*,std::size_t>> segments(1);
        segments[0].first = (void*)(s.data());
        segments[0].second = s.size();
        tl::bulk local = myEngine.expose(segments, tl::bulk_mode::write_only);
        b.on(ep) >> local;
        std::cout << "Server received bulk: ";
        std::cout << s;
        std::cout << std::endl;
    };

myEngine.define("do_rdma",f).disable_response();
```

Stream operators are defined to represent “push” and “pull” operations



## Bulk data transfers (cont'd)

Transferring subsegments of bulk regions

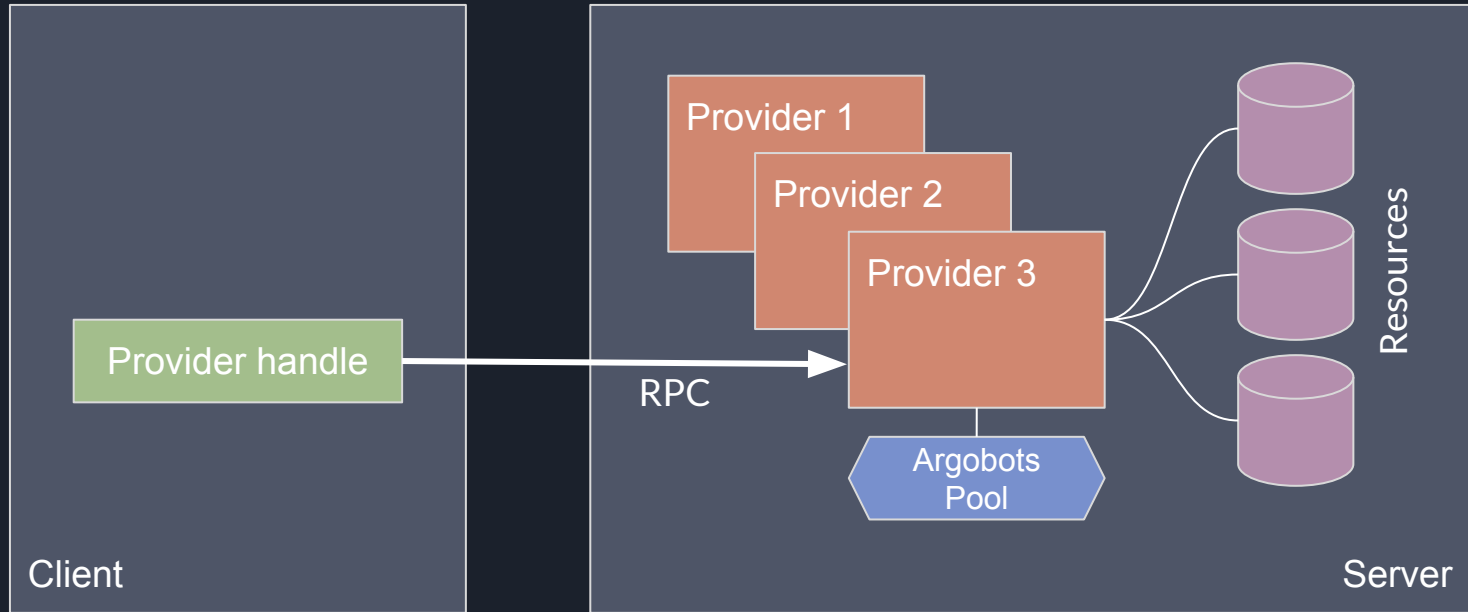
```
myRemoteBulk(3,45).on(myRemoteProcess) >> myLocalBulk(13,45);
```



Working in terms of  
providers



# What is a provider?



# Working in terms of providers

```
#include <iostream>
#include <thallium.hpp>
#include <thallium/serialization/stl/string.hpp>

namespace tl = thallium;

class my_sum_provider : public tl::provider<my_sum_provider> {

private:

    void sum(const tl::request& req, int x, int y) {
        std::cout << "Computing " << x << "+" << y << std::endl;
        req.respond(x+y);
    }

    int print(const std::string& word) {
        std::cout << "Printing " << word << std::endl;
        return word.size();
    }
}
```

Providers must inherit  
from `tl::provider`

You may omit the request  
argument if it's not used

By default, the non-void return value is  
treated as a response

## Working in terms of providers (cont'd)

```
public:

my_sum_provider(tl::engine& e, uint16_t provider_id=1)
: tl::provider<my_sum_provider>(e, provider_id) {
    define("sum", &my_sum_provider::sum);
    define("print", &my_sum_provider::print, tl::ignore_return_value());
}

~my_sum_provider() {
    wait_for_finalize();
}

};
```

## Working in terms of providers (cont'd)

```
public:
    my_sum_provider(tl::engine& e, uint16_t id) : tl::provider<my_sum_provider>(e, id) {
        define("sum", &my_sum_provider::sum);
        define("print", &my_sum_provider::print, tl::ignore_return_value());
    }

    ~my_sum_provider() {
        wait_for_finalize();
    }
};
```

Call define from the based class

Use this to indicate that although `print()` returns an `int`, we don't want to send this `int` back as a response

## Working in terms of providers (cont'd)

```
int main(int argc, char** argv) {  
  
    uint16_t provider_id = 22;  
    tl::engine myEngine("tcp", THALLIUM_SERVER_MODE);  
    std::cout << "Server running at address " << myEngine.self()  
                << " with provider id " << provider_id << std::endl;  
    my_sum_provider myProvider(myEngine, provider_id);  
  
    return 0;  
}
```

# Working in terms of providers (cont'd)

```
tl::engine myEngine("tcp", THALLIUM_CLIENT_MODE);
tl::remote_procedure sum = myEngine.define("sum");
tl::remote_procedure print = myEngine.define("print").disable_response();
tl::endpoint server = myEngine.lookup(argv[1]);
uint16_t provider_id = atoi(argv[2]);
tl::provider_handle ph(server, provider_id);
int ret = sum.on(ph)(42,63);
std::cout << "(sum) Server answered " << ret << std::endl;
std::string name("Matthieu");
print.on(ph)(name);
```

RPCs on clients should still be defined in the engine directly

A provider handle consists of an endpoint and a provider id

Call the RPC on a specific provider handle

# Properly finalizing providers







# Properly finalizing providers

```
class my_sum_provider : public tl::provider<my_sum_provider> {
    ...
    tl::remote_procedure m_sum;
    ...
    my_sum_provider(tl::engine& e, uint16_t provider_id=1)
    : tl::provider<my_sum_provider>(e, provider_id)
    {
        // keep the RPCs in remote_procedure objects so we can deregister them.
        m_sum(define("sum", &my_sum_provider::sum))
    }
    {
        // setup a finalization callback for this provider
        get_engine().push_finalize_callback(this, [p=this]() { delete p; });
    }
public:
    // this factory method and the private constructor prevent users
    // from putting an instance of my_sum_provider on the stack.
    static my_sum_provider* create(tl::engine& e, uint16_t provider_id=1) {
        return new my_sum_provider(e, provider_id);
    }
}
```



# Properly finalizing providers

```
~my_sum_provider() {
    m_sum.deregister();
    get_engine().pop_finalize_callback(this);
}

};

int main(int argc, char** argv) {
    tl::engine myEngine("tcp", THALLIUM_SERVER_MODE);
    myEngine.enable_remote_shutdown();

    my_sum_provider* myProvider22 = my_sum_provider::create(myEngine, 22);
    my_sum_provider* myProvider23 = my_sum_provider::create(myEngine, 23);

    return 0;
}
```

# Argobots wrappers





# Argobots wrappers

- barrier
  - condition\_variable
  - eventual
  - future
  - mutex
  - pool
  - recursive\_mutex
  - rwlock
  - scheduler
  - task
  - thread
  - timer
  - xstream
  - xstream\_barrier
- ABT\_barrier
  - ABT\_cond
  - ABT\_eventual
  - ABT\_future
  - ABT\_mutex
  - ABT\_pool
  - ABT\_recursive\_mutex
  - ABT\_rwlock
  - ABT\_scheduler
  - ABT\_task
  - ABT\_thread
  - ABT\_timer
  - ABT\_xstream
  - ABT\_xstream\_barrier

# Common pitfalls





# Common pitfalls

- Blocking MPI primitives (`MPI_Recv`, `MPI_Barrier`, etc.) don't yield to the progress loop. Use a dedicated progress thread, or a call MPI functions from a dedicated ES.
- Other threading frameworks (pthread, OpenMP, etc.) don't play well with Argobots threads. Mix with caution.
- Don't use `std::mutex`, remember to use `thallium::mutex`.
- Don't use `std::thread`, or `std::async`. Use corresponding Argobots equivalents.
- All the Argobots and Mercury objects should be released before margo/thallium is finalized.