



# Mochi affinity and threading

Mochi Bootcamp  
September 24-26, 2019



Argonne  
NATIONAL  
LABORATORY

# Process affinity





# Not all cores are created equal

Computer architectures are increasingly complex, particularly the nodes we typically see in HPC systems:

- Multi-core, multi-socket
- Numa nodes
- Multiple NICs, GPUs
- ...

To make most efficient use of these systems, it is important to take note of the locality of these devices and to allocate resources accordingly

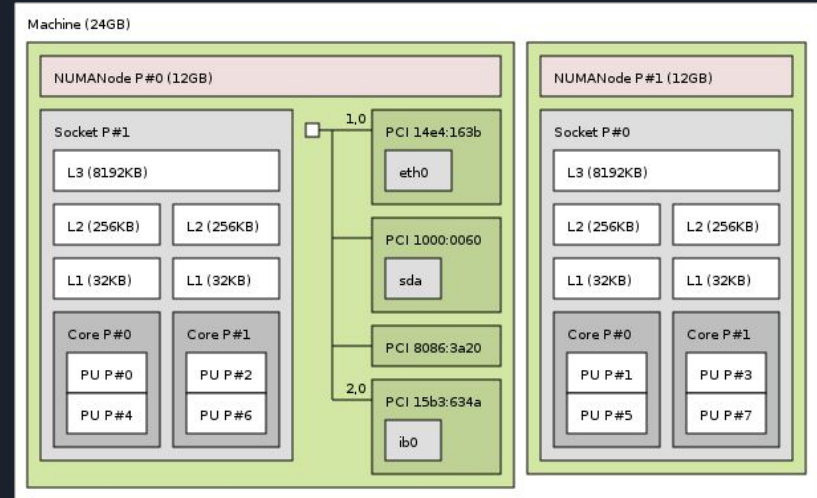
This problem is complicated by multi-threaded applications or by multiple applications/services sharing a node

# How can we learn about our target system?

## *hwloc*

- Provides a portable abstraction of the hierarchical topology of modern computer architectures

Using *hwloc*'s 'lstopo' command, we can generate graphical representations of our system architecture:





# Controlling processor affinity for Mochi apps

Now that you've learned more about your architecture, how can you actually take advantage of it?

## *Numactl*

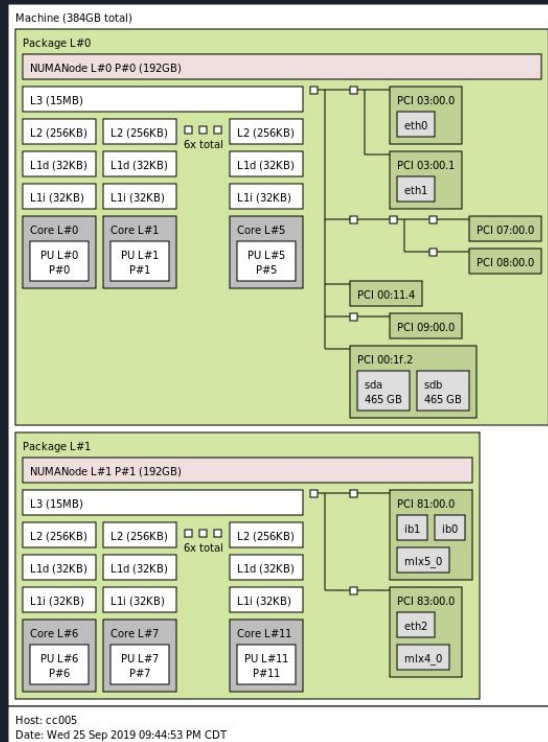
- Allows control of the NUMA scheduling and memory policies for a given executable
- Can specify which cores/sockets to run processes on, as well as which NUMA memory domain these processes allocate memory from
  - `--cpunodebind=sockets, -N sockets` : only execute on sockets
  - `--physcpubind=cpus, -C cpus` : only execute on cpus
  - `--membind=sockets, -m sockets` : only allocate memory from sockets

# hwloc/numactl example

Cooley Linux cluster @ ALCF, which uses an IB network

We prefer to pin a Mochi service on socket 1, since this socket shares locality with the IB controller for this node

➤ `numactl -N 1 -m 1 <executable>`



# Threading





# Threading

Getting comfortable with Argobots threading is critical to achieving desired performance under high-concurrency

Keep in mind some key Argobots terminology:

- *Execution stream (ES)* - sequential execution streams, essentially an OS thread
- *User-level threads (ULTs)* - an execution unit associated with a specific function
  - Scheduled on an ES, must yield to allow other ULTs execute on the ES
- *Scheduler* - Chooses what to execute from one or more associated pools
- *Pools* - set of schedulable work units for 1 or more ES

```
ABT_pool pool;  
ABT_pool_create_basic(ABT_POOL_FIFO_WAIT, ABT_POOL_ACCESS_MPMC,  
    ABT_TRUE, &pool);
```



# Threading

Getting comfortable with Argobots threading is critical to achieving desired performance under high-concurrency

Keep in mind some key Argobots terminology:

- *Execution stream (ES)* - sequential execution streams, essentially an OS thread
- *User-level threads (ULTs)* - an execution unit associated with a specific function
  - Scheduled on an ES, must yield to allow other ULTs to execute on the same ES
- *Pools* - set of units that can be produced and consumed on any ES
  - First-in, first-out pool with ability to wait gracefully
  - Units in pool can be produced on any ES and consumed on any ES

Automatically free pool

```
ABT_pool pool;  
ABT_pool_create_basic(ABT_POOL_FIFO_WAIT, ABT_POOL_ACCESS_MPMC,  
ABT_TRUE, &pool);
```



# Threading

Getting comfortable with Argobots threading is critical to achieving desired performance under high-concurrency

Keep in mind some key Argobots terminology:

- *Execution stream (ES)* - sequential execution streams, essentially an OS thread
- *User-level threads (ULTs)* - an execution unit associated with a specific function
  - Scheduled on an ES, must yield to allow other ULTs execute on the ES
- *Scheduler* - Chooses what to execute from one or more associated pools
- *Pools* - set of schedulable work units for 1 or more ES

```
...
```

```
ABT_sched sched;
```

```
ABT_sched_create_basic(ABT_SCHED_BASIC_WAIT, 1, &pool,  
    ABT_SCHED_CONFIG_NULL, &sched);
```

# Threading

Getting comfortable with Argobots threading is critical to achieving desired performance under high-concurrency

Keep in mind some key Argobots terminology:

- *Execution stream (ES)* - sequential execution streams, essentially an OS thread
- *User-level threads (ULTs)* - an execution unit associated with a specific function
  - Scheduled on an ES, must yield to allow other ULTs execute on the ES
- *Scheduler* - Chooses what to execute from one or more associated pools
- *Pools* - set of scheduled

Scheduler with ability to wait gracefully

1 or more pools to schedule work from

```
...  
ABT_sched sched;  
ABT_sched_create_basic(ABT_SCHED_BASIC_WAIT, 1, &pool,  
    ABT_SCHED_CONFIG_NULL, &sched);
```



# Threading

Getting comfortable with Argobots threading is critical to achieving desired performance under high-concurrency

Keep in mind some key Argobots terminology:

- *Execution stream (ES)* - sequential execution streams, essentially an OS thread
- *User-level threads (ULTs)* - an execution unit associated with a specific function
  - Scheduled on an ES, must yield to allow other ULTs execute on the ES
- *Scheduler* - Chooses what to execute from one or more associated pools
- *Pools* - set of schedulable work units for 1 or more ES

...

```
ABT_xstream xstream;  
ABT_sched_create_basic(sched, &xstream);
```



# Threading

Getting comfortable with Argobots threading is critical to achieving desired performance under high-concurrency

Keep in mind some key Argobots terminology:

- *Execution stream (ES)* - sequential execution streams, essentially an OS thread
- *User-level threads (ULTs)* - an execution unit associated with a specific function
  - Scheduled on an ES, must yield to allow other ULTs execute on the ES
- *Scheduler* - Chooses what to execute from one or more associated pools
- *Pools* - set of schedulable work units for 1 or more ES

...

```
ABT_xstream xstream;  
ABT_sched_create_basic(sched, &xstream);
```

Create an xstream associated  
with the given scheduler



# Providing xstreams/pools for Margo

At Margo init time, we have the opportunity to specify a couple of threading options:

Using regular `margo_init`:

- *use\_progress\_thread* - boolean value, 1 to use dedicated progress, 0 to use calling thread
- *rpc\_thread\_count* - number of ESs to allocate for RPC handlers, 0 to use calling thread, -1 to use progress thread

```
margo_instance_id margo_init_opt(const char *addr_str, int mode,  
int use_progress_thread, int rpc_thread_count);
```



# Providing xstreams/pools for Margo

At Margo init time, we have the opportunity to specify a couple of threading options:

Using regular `margo_init_pool`:

- *progress\_pool* - ABT\_pool to use for the progress thread
- *handler\_pool* - ABT\_pool to use for running RPC handlers

```
margo_instance_id margo_init_pool(ABT_pool progress_pool, ABT_pool handler_pool,  
hg_context_t *hg_context);
```

Note you need to also pass in an HG context, rather than an address string. This call is meant to provide caller most control over Margo init



# Providing xstreams/pools for Margo

At Margo RPC registration time, we can override the default handler pool we have specified at init time:

- Last argument is an ABT\_pool to use for executing handlers for the RPC type being registered

...

```
MARGO_REGISTER_PROVIDER(mid, "operation_name", void, void,  
    operation_ult, provider_id, pool);
```